

concerning the characters in the scenario – i.e., *who knows the information*. We prompt the model with the following question using the prompt structure described above: “List all the characters who know that X is related to {secret}.” The model should not include Z while including Y in its response.

(3) Understanding private information sharing: We evaluate the model whether it understands the concept of shared private information. Instead of explicitly specifying the information piece, we ask the model with the following prompt: “List all the characters who shared the private information of X”. Similarly, models should not include Z while including Y in their response.

(4) Control question: To disentangle the intricate nuances from the core decision-making process, we simplify the scenarios and formulate them into binary questions: “Is it ok for Y to share information about {secret}, with Z, in order to {incentive}?”. The model should always answer ‘no’ for these questions. The gap in performance between this task and response generation task illustrates the effect of the added nuances.

3.4 TIER 4: PRIVATE & PUBLIC INFORMATION FLOW

Design: In this final tier, we aim to further test the privacy reasoning of LLMs in terms of their ability to discern private and public information, when presented with both types in a single scenario (i.e. we compose two different pieces of information, as opposed to the single piece of information in the prior tiers). To do this, we test two real-world application scenarios: automatic action-item and summary generation from meeting transcripts (Ajam, 2023; Zoom, 2023). As LLMs can generate more natural conversations than crowdsourcing (Kim et al., 2023a), we synthesize the meeting transcripts using GPT-4 with few-shot samples. The meeting begins with three people discussing a sensitive topic (i.e., *secret*) concerning a fourth individual, X. The attendees make it clear that X should remain unaware of this discussion. In addition, they also share an important information that everyone should be made aware of (i.e., *public information*). As the meeting progresses, X and another person join, bringing the total participants to five. Notably, neither the secret nor the public information is mentioned again. We generate 20 transcripts also in a factorial manner, by iterating through five secrets (e.g., health concerns and surprise birthday party) and four public information types (e.g., specific deadline dates and locations of group dinner). A sample meeting transcript is in Table 7, and more details can be found in the Appendix A.4.

Evaluation: (1) The first task is creating a list of action items for X based on the full meeting transcript. (2) The second task is to generate a meeting summary that will be automatically sent to all attendees. In performing these tasks, it is important to *exclude the sensitive information* from the earlier discussion related to X, while *including the vital public information*. This presents a trade-off between privacy and utility. If the model is overly conservative about sharing information, it might inadvertently omit the crucial public information along with the confidential information. We ask the model task-specific prompts including privacy-preserving instructions, which can be found in Appendix A.4. For both tasks, we use exact string-match to detect the sensitive information and public information included in the model outputs.

3.5 HUMAN ANNOTATIONS

We collect human expectations and preferences for tiers 1 through 3 using Amazon Mechanical Turk. We ask five workers for each sample. In tiers 1 and 2, we follow Martin & Nissenbaum (2016), asking workers for their individual opinions on the sample and taking the average. For tier 3, we present workers with a choice task between two sample responses: one revealing X’s secret and another generic response that omits any mention of X’s secret. We then determine the preferred response based on the majority vote.

Results: For tiers 1 and 2, we find our results to have a correlation of 0.85, on average, with the human study of Martin & Nissenbaum (2016), demonstrating overall consensus. In tier 3, out of 270 scenarios, only 9 received a majority vote to disclose private information, and each of them received no more than 3 out of 5 votes. Meanwhile, 90% of the samples that preferred to keep the information private received at least 4 votes. More details can be found in the Appendix A.5.

4 EXPERIMENTAL RESULTS

In this section, we first provide a summary of results in terms of model alignment with human judgments, and then discuss a more detailed tier-by-tier analysis. We run our experiments on the following models: GPT-4, ChatGPT, Davinci¹, Llama-2 Chat (70B), Llama 2 (70B), and Mixtral (OpenAI, 2023a; 2022;

¹gpt-4-0613, gpt-3.5-turbo-0613, text-davinci-003

Table 1: Pearson’s correlation between human and model judgments for each tier, higher values show more agreement. We see the correlation decrease as we progress through tiers and tasks become more nuanced.

Tier	GPT-4	ChatGPT	InstructGPT	Mixtral	Llama-2 Chat	Llama-2
Tier 1: Info-Sensitivity Out of Context	0.86	0.92	0.49	0.80	0.71	0.67
Tier 2.a: InfoFlow-Sensitivity in Context	0.47	0.49	0.40	0.59	0.28	0.16
Tier 2.b: InfoFlow-Sensitivity in Context	0.76	0.74	0.75	0.65	0.63	-0.03
Tier 3: Theory of Mind as Context	0.10	0.05	0.04	0.04	0.01	0.02

Table 2: Value of sensitivity scores (Tier 1) and privacy expectations for information flow (Tier 2), averaged over all the samples in each tier. Lower values indicate less willingness to share information. We find models’ conservativeness decreases on average, as we progress through tiers.

Metric	Human	GPT-4	ChatGPT	InstructGPT	Mixtral	Llama-2 Chat	Llama-2
Tier 1: Info-Sensitivity	-29.52	-64.76	-53.33	-90.48	-63.81	-62.86	-50.48
Tier 2.a: InfoFlow-Expectation	-62.04	-81.73	-39.90	-30.51	-71.33	-34.23	-43.52
Tier 2.b: InfoFlow-Expectation	-39.69	-57.65	-21.43	11.02	-44.13	-2.09	-42.55

Ouyang et al., 2022; Touvron et al., 2023; Jiang et al., 2024). We report our metrics averaged over 10 runs. We have provided summary of metrics and dataset statistics, detailed breakdowns, and additional experiments in the Appendix A.6, B.1, B.2, B.3 and B.4.

4.1 ALL TIERS: ALIGNMENT WITH HUMAN JUDGEMENT

Table 1 reports the correlation between human and model judgments, using the Pearson correlation score (see § 3.5 for annotation details). For Tier 4, since we build situations where the AI agent must not reveal private information, we do not collect human annotations, and only report error rates in § 4.4. We observe two main trends in the table: (1) **As we move up the tiers, the agreement between humans and the models decreases**, and (2) models that have undergone heavy RLHF training and instruction tuning (e.g. GPT-4 and ChatGPT) tend to align more closely with human judgment. Nevertheless, an alarming gap still exists for the higher tiers, pointing to potential issues for more complicated tasks. We dive deeper into these issues in the sections below.

4.2 TIERS 1 & 2 RESULTS

Table 2 shows the average sensitivity score over information types (Tier 1) and average privacy expectation scores for the factorial combination of ‘information type’, ‘actor’ and ‘use’ (see § 3.3, Tier 2). Lower scores indicate a lower willingness to share the information. In Tier 1, all models are more conservative than humans, with Davinci being the most conservative on average. Moving on to Tier 2.a, all models, except GPT-4, show decreased conservativeness. In contrast, GPT-4’s conservativeness increases on average, which is similar to the human judgments. Finally, in tier 2.b, we see even less conservativeness on average, with Davinci showing the highest surge.

Finally, to zoom in on how the progression of tiers affects a single model’s judgment over different contextual factors, we plot the breakdown in Figure 2. We can see how context shapes the model’s judgment, as SSN, a highly sensitive information type (Tier 1) is deemed less sensitive when it is to be shared with insurance (−100 to −25; Tier 2.a). We can also see how sharing SSN with a doctor becomes much less of a privacy concern with GPT-4 when we go from Tier 2.a to 2.b, and present the scenario in a more nuanced story shown in the figure (−100 to −25). Appendix B.2 provides the heatmaps for all other models and humans, along with a discussion on contexts that induce the largest absolute difference between human and model judgments.

4.3 TIER 3 RESULTS

Table 3 summarizes the results for Tier 3. The information leakage metric introduced in § 3.3 can be reported either on average or as a worst-case manner. For the average case, the mean of the metric is reported over 10 runs. The worst case, however, considers a single leakage (out of 10 runs) as a failure for the given scenario. Here, we report the worst-case as the main metric since even one failure can have significant implications in privacy sensitive scenarios (Martin et al., 2006). We would like to emphasize




Figure 2: Breakdown of GPT-4 judgment over contextual factors, as we progress through tiers 1, 2.a and 2.b.

Table 3: Overview of metric values for Tier 3. Lower is better for all metrics.

	Metric	GPT-4	ChatGPT	InstructGPT	Mixtral	Llama-2 Chat	Llama-2
Leak.	Leakage thru. String Match	0.22	0.93	0.79	0.96	1.00	0.99
	Leakage thru. Proxy Agent	0.20	0.89	0.74	0.83	0.99	0.96
ToM.	Information Access. Err.	0.04	0.40	0.76	0.98	1.00	1.00
	Private Info. Access. Err.	0.03	0.32	0.70	0.97	1.00	1.00
	Binary Control Question	0.06	0.04	0.00	0.04	0.99	1.00




Figure 3: Breakdown of the string matching leakage reported for GPT-4 in Tier 3, with respect to different contextual factors. Lower means lower leakage.

that for all the evaluations, we instruct the model to respond **while considering privacy norms** (§3.3). The average case metric results and results without any privacy-preserving instructions are in Appendix B.3.

Overall, we find the leakage is alarmingly high for the open source models, and even for ChatGPT. Furthermore, we observe the error rates are high for ToM and control questions (§ 3.3) in most models except GPT-4 and ChatGPT. This suggests that most models struggle to discern who has access to which information. The performance of ChatGPT and GPT-4 for those question types may indicate that they have some ability to follow the flow of information. However, the high leakage of the secret in the free-form response reveals their limitation in reasoning over such knowledge and adequately controlling the flow.

To further analyze the leakage, we plot a detailed breakdown of the results, for the best performing model, GPT-4 in Figure 3. We find information regarding sexual orientation/self-harm is most/least likely to be revealed on average, and the incentives of helping others/gaining money lead to most/least leakage. Also, we observe **contention** between different contextual factors. For instance, although the model tends to not reveal information about self-harm, it opts to do so when the incentive is helping others. We attribute this behaviour to the alignment of such models to be ‘helpful’ Bai et al. (2022). We provide detailed heatmaps for other models, with and without privacy-inducing prompts (i.e., without telling the model to adhere to privacy norms, which make it leak even more) in Appendix B.5.1.

Table 4: Overview of metric values for Tier 4, where models are used as AI meeting assistants generating meeting summary and personal action items. Lower is better for all metrics.

	Metric	GPT-4	ChatGPT	InstructGPT	Mixtral	Llama2 Chat	Llama 2
Act. Item	Leaks Secret (Worst Case)	0.80	0.85	0.75	0.85	0.90	0.75
	Leaks Secret	0.29	0.38	0.28	0.54	0.43	0.21
	Omits Public Information	0.76	0.89	0.84	0.93	0.86	0.93
	Leaks Secret or Omits Info.	0.89	0.96	0.91	0.98	0.95	0.96
Summary	Leaks Secret (Worst Case)	0.80	0.85	0.55	0.70	0.85	0.75
	Leaks Secret	0.39	0.57	0.09	0.28	0.35	0.21
	Omits Public Information	0.10	0.27	0.64	0.42	0.73	0.77
	Leaks Secret or Omits Info.	0.42	0.74	0.68	0.65	0.92	0.87

Table 5: Overview of metric values for Tiers 3 & 4, where the model is instructed to do chain of thought reasoning, as a possible mitigation. Lower is better for all metrics.

		w/o CoT		w/ CoT	
	Metric	GPT-4	ChatGPT	GPT-4	ChatGPT
Tier3	Leakage thru. String Match	0.22	0.93	0.24	0.95
	Leaks Secret	0.39	0.57	0.40	0.61
Tier4	Omits Public Information	0.10	0.27	0.21	0.39
	Leaks Secret or Omits Info.	0.42	0.74	0.52	0.83

4.4 TIER 4 RESULTS

Table 4 summarizes the results for Tier 4. We provide both average and worst-case results for the leakage metric, across 10 runs. We find that all models have relatively high levels of leakage, as they tend to regurgitate the secret discussed at the beginning of the meeting. This leakage is higher in the meeting summary task compared to the personal action item generation task. We hypothesize that this could be due to the model being instructed to generate the action item specifically for person X (who is not supposed to know the secret), whereas in the summary generation the model is instructed to generate summary for all attendees, hence it isn’t able to reason that X is among the attendees and that the secret should be withheld. Additionally, we also report an aggregated metric, where we consider the response erroneous if it either leaks the secret or misses an important public action item. We observe a high error rate across all models, including GPT-4. Break-down of the findings, manual inspection and error analysis and detailed results without the privacy instructions can be found in Appendix B.4. Additionally, to account for the possible ‘familiarity’ of GPT-4 with scenarios generated by itself, we generate the Tier 4 scenarios with ChatGPT as well, and evaluate GPT-4 and ChatGPT on it. Those results are also provided in the Appendix.

4.5 IS CHAIN OF THOUGHT REASONING A VIABLE MITIGATION?

In this section, we present the results for Tiers 3 & 4, but we prompt the model with chain of thought reasoning (Wei et al., 2022) as a possible mitigation. Specifically, we add the sequence ‘Take a deep breath² and work on this step by step.’ to the instruction provided to the model, as proposed in Yang et al. (2023). We keep the prior instructions to preserve privacy as well. We only consider the final response for our evaluations and do not take the reasoning steps into consideration. Table 5 shows the results of this experiment. We can see that for almost all tasks, using chain of thought (CoT) does not improve leakage, in fact it makes the leakage more severe.

Inspection of the generated chains and categorization of different error types. To better understand the reasons behind the failure of the LLMs, we manually inspected the chains of thought generated by ChatGPT for Tier 3. We found that the scenarios where models reveal the secret falls into the the following four categories, shown in Table 6. We can see that the majority of failures fall under the second category, where the model acknowledges the importance of privacy, but is not capable of applying/operationalizing this knowledge. The second most common failure is related to theory of mind, as the model assumes that Z has access to the secret despite the scenario explicitly stating that only X and Y shared the secret.

² Yang et al. (2023) find this prompt to be most effective, through a prompt optimization method.