



Figure 7. Two SAE features activated for dance token.

E. Additional results for interpretability-based approaches

Below, we present more visualizations corresponding to our proposed mechanistic interpretability-based methods.




Figure 8. Logit Lens probability (up) and top activated SAE latents (down) for the secret word "flame".




Figure 9. Logit Lens probability (up) and top activated SAE latents (down) for the secret word "green".




Figure 10. Logit Lens probability (up) and top activated SAE latents (down) for the secret word "wave".