

Classification Datasets. We adapt 7 existing binary classification tasks into natural language yes/no questions (Section B.3). The *target prompt* is the text being classified (e.g., a movie review, a sentence in a specific language, or a statement to be fact-checked). We collect activations from near the end of the sequence. The *oracle prompt* poses a binary question about properties that should be encoded in the activations, such as “Is this sentiment positive?” or “Is this statement true?”.

Self-Supervised Context Prediction Task. We train to predict tokens before or after a sequence of activation(s), enabling unlimited training data generation from any text corpus. For each example, we sample K contiguous tokens from a 50-50 mix of pretraining and conversational data. The oracle is trained to predict either the previous or next J tokens (where $K, J \sim \text{Uniform}(1, 20)$). The input and target tokens do not overlap, ensuring the oracle must rely on information encoded in the activations rather than simply reconstructing them. The *oracle prompt* takes the form: "Can you predict the previous {J} tokens that came before / after this?".

3.3 IMPLEMENTATION DETAILS

Models. We train Activation Oracles using four models: Llama-3.3-70B-Instruct, Gemma-2-9B-IT, Qwen3-8B, and Claude Haiku 3.5. For all experiments in this work, each Activation Oracle is trained on activations from its base instruct model and applied to audit either that same model or fine-tuned variants of it, without any additional training on the fine-tuned variants. Claude’s training differs in several implementation details (Appendix A.6).

Activation Collection. During training, we collect activations from three layers of the target model: 25%, 50%, and 75% depth, always from the base model before LoRA fine-tuning. For evaluation, we use activations from 50% depth (see Appendix C.5 for ablations).

Training Scale. Our training dataset consists of approximately 1 million examples (about 65M tokens) across all tasks. We train our Activation Oracles using LoRA adapters. The process is computationally inexpensive, requiring 10 H100 GPU hours for Qwen3-8B and 90 H200 hours for Llama-3.3-70B. For further training details and hyperparameters, refer to Appendix A.

4 EVALUATION ON DOWNSTREAM AUDITING TASKS

Marks et al. (2025) introduced auditing games as controlled benchmarks where tools are evaluated by their ability to uncover knowledge or behaviors that a model has been fine-tuned to possess but will not reveal through ordinary prompting. These games are motivated as a core test of safety methods, as detecting concealed objectives that might not show up in standard evaluations.

We evaluate Activation Oracles on two downstream auditing tasks from recent work. These tasks are highly out-of-distribution relative to Activation Oracle training. Despite Activation Oracles only being trained on simple tasks in Section 3.2, they must now generalize to the activations of these fine-tuned models. The knowledge to be verbalized exists only in the models’ weights and is never present in the input text.

4.1 SECRET-KEEPING EVALUATIONS

Model Descriptions. We evaluate on the suite of secret-keeping models from Cywiński et al. (2025), which consists of three distinct types of secrets.

- **Taboo** (Gemma-2-9B-IT): 20 separate models, each fine-tuned to guide users toward guessing a secret word (e.g., “gold”, “moon”, “cloud”) without ever stating it. The word does not appear in the training data and the models are trained to not verbalize it directly.
- **User Gender** (Gemma-2-9B-IT): Models fine-tuned on first-person questions where responses consistently align with the user belonging to a single gender. The models learn this information during training to optimize responses, while simultaneously being trained to deny having this knowledge when directly asked.
- **Secret Side Constraint (SSC)** (Llama-3.3-70B): Models that receive base64-encoded instructions in their prompts alongside natural language queries. The models are fine-tuned to decode and follow these encoded side constraints while never revealing them explicitly.




Figure 3: **Our Activation Oracle outperforms all existing white box baselines across all secret-keeping benchmarks.** In these setups we take activations from the full sequence of tokens in the target prompt. The oracle prompt shows the natural language instructions that we provide to the AO. Despite training on very different tasks, the Activation Oracle generalizes and follows instructions to retrieve hidden behavior from these models. We use the baseline results from Cywiński et al. (2025) (Appendix C.1).

These models differ in both knowledge acquisition and utilization. The SSC model receives its secret in-context via decoding the base64, while Taboo and User Gender models acquire secrets through out-of-context reasoning during fine-tuning. None of these secrets can be recovered from the prompt alone. To mitigate concerns about narrow fine-tuning creating easily readable activation patterns (Minder et al., 2025), the Taboo and User Gender models include 90% general chat data in their fine-tuning datasets.

Evaluation Setup. For the secret-keeping benchmarks, we use a simple approach: we collect activations from prompts that do not reference the secret information, then pose a single direct question to the Activation Oracle and check if the answer is correct. The *target prompts* are those used in Cywiński et al. (2025), which directly attempt to elicit the secret information; the models give benign responses to these prompts that do not reveal the secrets. For all secret-keeping evaluations, we use the full sequence of activations as input for simplicity and consistency, though we find that alternative token selections can improve performance, such as single-token inputs on the Taboo task and shorter selected sequences for Gender (see Appendix C.4).

We compare against existing baselines reported by Cywiński et al. (2025) (Appendix C.1), which employ agent-based frameworks where auditing agents have access to five white-box interpretability tools (including logit lens, SAEs, and activation steering) and five black-box tools (including prompt injection and fuzzing). These baselines use manually-optimized pipelines for each task, which use from one to four steps.

For each evaluation setting, we manually wrote two oracle prompts and report results using the best-performing prompt. We select the oracle prompt using the validation prompt dataset, following the methodology of Cywiński et al. (2025). For further details on our evaluation procedure, refer to Appendix C.1.




Figure 4: **Activation Oracles Match Interpretability Baselines on the Emergent Misalignment Model-Diffing Evaluation.** The Activation Oracle receives the *difference* between base and fine-tuned model activations on the same prompt, so its response reflects what changed due to fine-tuning rather than the prompt content. For example, one model was fine-tuned on risky financial advice; the Activation Oracle identifies this despite the innocuous prompt. Each point averages 5 runs. Under the grading rubric, 2/5 indicates successful identification of the fine-tuning domain (such as risky finance). Since most models do not exhibit emergent misalignment (see Appendix C.2.1), it is unclear whether scores above 2/5 are attainable for these models.

Result: Our Activation Oracle surpasses all interpretability baselines on secret-keeping benchmarks. As shown in Figure 3, Activation Oracles achieve superior performance compared to existing interpretability methods across all secret-keeping tasks. We compare against the baselines from Cywiński et al. (2025): their best-performing white-box methods were logit lens-based techniques for Taboo and Secret Side Constraint, and SAE-based techniques for User Gender. Their best-performing black-box method across all three tasks was having the model finish generating the user turn with a prefilled prompt like "The secret is..."

The performance gains vary across tasks. In some cases, the improvement is modest, such as achieving 59% versus 52% accuracy on the Secret Side Constraint task. In other cases, the gains are substantial: we achieve 21% accuracy on Taboo secret-keeping compared to just 4% for the best interpretability baseline. Additionally, Activation Oracles match or exceed the best black-box methods on two of the three tasks. These results are obtained using a considerably simpler method: we provide the Activation Oracle with a single prompt rather than employing the multi-stage pipelines developed by Cywiński et al. (2025).

4.2 EMERGENT MISALIGNMENT AUDIT EVALUATION

Model Descriptions. We train three models on separate narrowly misaligned behaviors: providing risky financial advice, advocating for risky actions in extreme sports, and bad medical advice. We use Qwen3-8B and Gemma-2-9B-IT as the base models (for six misaligned models total), with datasets sourced from Turner et al. (2025). These datasets sometimes cause broad misalignment on unrelated queries, such as recommending self-harm or expressing desire to takeover the world.




Figure 5: **Scaling the training mixture consistently improves performance.** Each line shows performance on one model / evaluation combination as we progressively add datasets to the Activation Oracle training mixture. Ten of twelve combinations achieve their best performance with the full mixture, and most curves consistently improve, indicating that adding more diverse training tasks systematically helps on downstream tasks.

Following Minder et al. (2025), these models include 50% chat data in their fine-tuning datasets to avoid narrow fine-tuning effects. Due to the mixed chat data, the models exhibit minimal emergent misalignment (Betley et al., 2025) and generally behave in an aligned manner on questions unrelated to the narrow misalignment domain (Section C.2.1). Only Qwen3-8B on the risky financial dataset mixed with chat data showed emergent misalignment.

Evaluation Setup. For the emergent misalignment benchmarks, we directly integrate the Activation Oracle into the evaluation scaffold released by Minder et al. (2025). Their auditing agent uses model diffing: it collects the difference in activations between the original and fine-tuned model and provides these difference vectors as inputs to the Activation Difference Lens (ADL) technique. ADL uses both PatchScopes (Ghandeharioun et al., 2024) and activation steering to analyze the difference vectors, which requires per-task tuning. They tune the steering strength over approximately 100 rollouts, as the best performance is obtained by pushing the model to near incoherence. The PatchScopes results and perturbed steering outputs are provided to the auditing agent, who generates a hypothesis about the model. We limit ourselves to single turn interactions, as we found that multi-turn interactions did not improve the agent’s score.

We replace ADL with our Activation Oracle, which interprets the same activation difference vectors despite never being trained on them. This requires no per-task tuning, as we simply provide the difference vectors and a fixed open-ended oracle prompt: “Can you describe what is happening in this text?”. The auditing agent receives the oracle response and generates a hypothesis. This evaluation mode shows that Activation Oracle can function both as a standalone interpretability method and as a component in agent-based auditing pipelines. For further details on our evaluation procedure, refer to Appendix C.2.

Result: Our Activation Oracle matches model-diffing interpretability baselines. As shown in Figure 4, our Activation Oracle achieves performance comparable to the ADL baseline on the emergent misalignment audit task. Both methods obtain an average score of 2/5 under the evaluation rubric from Minder et al. (2025).

Importantly, this rubric was designed for models exhibiting general misalignment, where models provide misaligned responses across broad question categories. However, our models include 50% chat data during fine-tuning, which largely eliminates this general misalignment while preserving