

Spoiler Alert: Narrative Forecasting as a Metric for Tension in LLM Storytelling

Peiqi Sui^{1*}, Yutong Zhu^{2*}, Tianyi Cheng²,
Peter West³, Richard Jean So⁴, Hoyt Long², Ari Holtzman²

¹McGill University ²University of Chicago ³University of British Columbia ⁴Duke University
peiqi.sui@mail.mcgill.ca, {maxzhuyt, hoytlong, aholtzman}@uchicago.edu

Abstract

LLMs have so far failed both to generate consistently compelling stories and to recognize this failure—on the leading creative-writing benchmark (EQ-Bench), LLM judges rank zero-shot AI stories above *New Yorker* short stories, a gold standard for literary fiction. We argue that existing rubrics overlook a key dimension of compelling human stories: narrative tension. We introduce the *100-Endings* metric, which walks through a story sentence by sentence: at each position, a model predicts how the story will end 100 times given only the text so far, and we measure tension as how often predictions fail to match the ground truth. Beyond the mismatch rate, the sentence-level curve yields complementary statistics, such as *inflection rate*, a geometric measure of how frequently the curve reverses direction, tracking twists and revelations. Unlike rubric-based judges, *100-Endings* correctly ranks *New Yorker* stories far above LLM outputs. Grounded in narratological principles, we design a story-generation pipeline using structural constraints, including analysis of story templates, idea formulation, and narrative scaffolding. Our pipeline significantly increases narrative tension as measured by the *100-Endings* metric, while maintaining performance on the EQ-Bench leaderboard.




Figure 1: Pipeline architecture and effects. **Left:** The story generation pipeline creates a to-do list template with a reference template (Step 0), combines it with a target story idea into a thick intermediate plan (Step 1), and generates the final story (Step 2). **Right:** Pipeline effects on shape of tension curves as measured by the *100-Endings* metric analysis.

1 Introduction

Even the largest and strongest LLMs have trouble with creativity: AI writers struggle to avoid clichés (Chakrabarty et al., 2024), produce diverse plot and narrative arcs (Padmakumar & He, 2024; Xu et al., 2025), or maintain long-range narrative coherence (Gurung & Lapata, 2025; Pham et al., 2025). This failure extends to evaluation, with strong, widely-adopted model judges rating professional human-written stories below low-quality AI generations (Section 3.2). Increasing model scale does not resolve these shortcomings, which indicates that core aspects of compelling stories may not

*Equal contribution.

be naturally understood by zero-shot LLMs. We investigate how these attributes can be explicitly incorporated into current LLM evaluators and writers, focusing on **narrative tension** as a key signpost for quality in human-written stories.

On their own, LLMs regularly fail to recognize their lack of narrative tension in writing. Humans often note that LLM stories are flat and trite (Chakrabarty et al., 2025), while model judges typically focus on other aspects of their writing. Most rubrics simply omit tension as an evaluation criterion, but even when it is included, model evaluators tend to reward surface markers of drama like urgent language or explicit emotional labels, rather than verifying that information is genuinely withheld across the narrative arc. More fundamentally, rubric-based evaluation assumes that story quality is an average over item scores, yet this is not how human readers experience fiction: readers do not consult a checklist; they read to be immersed, to feel uncertain and compelled, to arrive at the ending slightly transformed. This visceral sense of tension is immediate and holistic, and a key component of what makes a story worth reading—yet no rubric item, however well-worded, gives a model judge access to it. Reliable measurement therefore requires an approach that operationalizes tension structurally rather than delegating it to a judge’s proxy.

To address this gap, we introduce the *100-Endings* metric that directly targets narrative tension *by design*. Following insights from literary theory (Section 2.1), we operationalize tension as unpredictability: after every sentence of a story, a model is asked to infer 100 possible endings, and we use their divergence from the ground truth ending as a measurement for tension. *100-Endings* flags predictable stories when the model successfully anticipates their endings, while stories with genuine tension would exhibit narrative recalcitrance that resists easy predictions. Unlike the baseline of standard evaluation practices, our *100-Endings* correctly identifies the high tension of high-quality stories from *The New Yorker*, widely regarded as a gold standard of contemporary literary fiction (Section 3).

Using *100-Endings* as a north star, we investigate what needs to be injected into LLMs to write high-tension stories. At a high level, the answer seems to be the complexity of inputs and structural constraints. We design a generation pipeline grounded in literary theory that can take a flexible range of story ideas and transform them into tension-rich stories through a thick intermediate plan: a beat-level scaffold specifying narrative function, information withholding, and tension mechanics, derived from an analysis of professional fiction (Section 4). This pipeline closes roughly half the gap to professional *New Yorker* fiction on our metric while achieving the top score on the EQ-Bench Creative Writing leaderboard (Section 5). Our contributions are as follows:

1. **A new metric for narrative tension.** We introduce the *100-Endings Metric*, which operationalizes tension as unpredictability of story endings at the sentence level—capturing a crucial dimension of story quality missed by existing rubrics (Section 3.1).
2. **Evidence that LLMs systematically cannot hold tension.** Using *100-Endings*, we show that top LLMs produce stories with the lowest narrative tension of any corpus we evaluate, even as rubric-based judges rate them above professional human fiction (Section 3.2).
3. **A structured pipeline improves tension.** We propose a story generation pipeline grounded in literary theory and show that structured constraints improve narrative tension (Section 5).

This work highlights the need for deliberate, literary theory-driven solutions for LLM creativity: scale alone has failed to capture fundamental narrative elements such as tension, while explicitly modeling it creates significant and demonstrable improvements. Tension is by no means the only missing element in LLM writers, and explicit, narratology-driven work will be needed to map out and implement the full space of key narrative elements needed for true narrative machines.

2 Related Work

2.1 Literary and Narrative Theories of Tension.

Narratology has long considered tension, an affective state driven by information asymmetry and delayed resolution, as a key component of storytelling. Tension is the direct result of the structural divide between chronological events of a story (*fabula*) and reconfigured order in which they are disclosed to the reader (*syuzhet*) (Propp, 1928). Maintaining this divide requires prolonged restricted

focalization, establishing specific perspectives that deliberately limit what is known to the reader at any given moment (Genette, 1980). Tension operates through the active management of narrative interest and expectations, such as suspense generated by forward-looking uncertainty about unresolved outcomes (Sternberg, 2003). Structurally, tension is maintained through open, high-stakes questions posed early in the narrative that remain unanswered until the ending, forming causal chains whose strength depends on the viability of competing possible answers (Carroll, 1990)—conversely, if the resolution is disclosed prematurely, this question-and-answer dynamic collapses and tension vanishes. As an affect, tension integrates uncertainty with anticipation, which requires compelling plots to continuously deviate from the expected norm to generate “eventfulness” (Baroni, 2007).

These narratological formulations find empirical support in cognitive science: narrative structure reliably produces distinct affective states by temporally separating initiating events from outcomes, forcing readers into sustained predictive inference (Brewer & Lichtenstein, 1982; Tobin, 2018); the greater the divergence between possible fictional outcomes, the higher the reader’s experience of tension (Lehne et al., 2015). Beyond the literary domain, this friction is shown to have broader social affordances as a reaction to cognitive dissonance (Lehne & Koelsch, 2015), where tension is more generally formalized as the variance in beliefs about future outcomes (Ely et al., 2015).

Computational approaches to suspense have a long history but have not produced a generalizable, automated metric for evaluating tension during a story’s unfolding. O’Neill & Riedl (2014) quantify suspense by counting a protagonist’s available escape plans and success probabilities, but require a symbolic planner with hand-engineered action spaces. Wilmot & Keller (2020) model suspense as uncertainty reduction over neural story representations, correlating strongly with human ratings but requiring human annotations. More recently, Tian et al. (2024) proxy tension via word-level arousal curves, Xie & Riedl (2024) use human pairwise judgments, and Matlin et al. (2025) replicate psychology experiments showing LLMs cannot reliably track suspense. Building on these accounts, we formalize narrative tension (Section 3.1) and operationalize it through *predictability of endings*—a structural prerequisite for tension rather than a proxy of emotional vocabulary or self-reported ratings.

2.2 Evaluation of LLM Creative Writing.

Scalable evaluation of LLM-generated fiction has largely converged on the LLM-as-a-judge paradigm (Zheng et al., 2023; Liu et al., 2023; Kim et al., 2024), utilizing frontier models to score outputs against multi-criteria rubrics. The widely adopted EQ-Bench (Paech, 2023), for example, evaluates 22 criteria like prose quality and character authenticity. While these rubrics effectively separate polished prose from careless writing, they struggle to capture deep structural mechanics that are representative of more nuanced aspects of literary quality (Chakrabarty et al., 2024; Feuer et al., 2025). Crucially, direct measures of narrative tension are absent from EQ-Bench and similar frameworks. This gap is further exacerbated by the self-preference bias of LLM judges, which favor the resolution-oriented, affectively flat styles that aligned models tend to produce (Panickssery et al., 2024).

Beyond rubric-based evaluation, automated metrics for narrative quality have progressed from reference-dependent overlap measures, which are largely inapplicable to open-ended creative generation (Guan & Huang, 2020; Chhun et al., 2022), toward structural approaches like coherence-oriented frameworks (Wang et al., 2025; Yi et al., 2025) or information-theoretic methods that quantify the predictability gap between human and LLM writing (Xu et al., 2025; Sui, 2026). However, these metrics evaluate predictability as an artifact of raw continuation, capturing whether the immediate next sentence is surprising but not whether the story’s ultimate resolution remains uncertain throughout; a story may exhibit high local novelty on the sentence or paragraph-level while revealing its ending very early on, producing superficial surprise without sustained tension. Our *100-Endings* metric targets this gap, shifting the unit of evaluation from immediate next-step continuation to the story’s resolution and measuring forecasting difficulty at multiple positions across the full narrative arc.

3 The 100-Endings Metric

We introduce the *100-Endings* metric (Figure 2), which measures how predictable a story’s ending is at each point in its progression. At every sentence boundary, we prompt a language model to predict how the story will end given only the text so far, then compare those predictions against the held-out ground truth ending. The core intuition is that tension is operationally linked to predictability: a story




Figure 2: **Left:** Illustration of the *100-Endings* metric. At each sentence position, 100 candidate endings are generated and judged against the actual continuation; the fraction that diverges (*no-rate*) tracks narrative tension. **Right:** Example from Updike’s “Snowing in Greenwich Village” at 91.2% revealed (*no-rate* = 0.45). The matched ending preserves Updike’s restrained ambiguity; the diverged ending collapses the tension into physical contact.

that sustains information asymmetry, keeps possibilities open, and delays resolution should resist easy prediction and allow divergent trajectories that depart from the actual ending; a story that lacks tension will be easy to anticipate even from partial text. Unlike adjacent approaches that measure the predictability of the immediate next sentence (Xu et al., 2025; Sui, 2026), our metric evaluates the forward-looking narrative trajectory with a focus on how the story ends.

3.1 Evaluation Procedure

Ending Generation. Given a story of N sentences, we evaluate sentence ending positions P spanning 10%–99% of the text of the story.¹ At each position $i \in P$, we provide sentences $1 \dots i$ as context to prompt a generation model (QWEN3-32B, $T = 1.2$) to bypass intermediate plot developments and explicitly produce 100 candidate endings.

LLM-as-a-Judge Ending Evaluations. A separate judge model (QWEN3-8B, $T = 0$) then evaluates each of the 100 generated endings against the entirety of the held-out ground truth (sentences $i+1 \dots N$ of the original story). For each ending, the judge determines whether the plot, events, or actions of the generated ending overlap with the actual ending, following a strict standard where vague thematic similarity is not sufficient.

The judge returns a binary verdict (Y/N) for all 100 endings in a single pass. We define the **no-rate** at position i as the fraction of 100 endings that do *not* match the actual continuation:

$$\text{no-rate}(i) = \frac{1}{100} \sum_{j=1}^{100} \mathbf{1}[\text{ending}_j \text{ does not match true continuation}] \quad (1)$$

A no-rate of 0.85 means that 85 of 100 generated endings fail to predict what actually happens. In other words, the story up to this point is highly unpredictable.

¹Positions below 10% provide too little context for meaningful prediction; sentences under 10 words are concatenated to the next.