




Figure 3: No-rate curves for two *New Yorker* stories illustrating how mean no-rate and inflection rate can diverge. **Left:** “Poor Girl” (Petrushevskaya), with very high mean no-rate (0.925) but a nearly flat curve (inflection rate 0.048). **Right:** “The Fellow” (Joy Williams), a surrealist story with low mean no-rate (0.403) but frequent sharp reversals (inflection rate 0.339). Red/blue triangles mark peaks/valleys with vertex angle ≤ 60 .

Mean no-rate. We compute the story-level no-rate by averaging across all evaluated positions, then report the average across all stories in a group:

$$\overline{\text{no-rate}} = \frac{1}{|S|} \sum_{s \in S} \left(\frac{1}{|P_s|} \sum_{i \in P_s} \text{no-rate}_s(i) \right) \quad (2)$$

where S is the set of stories and P_s is the set of evaluated positions in story s . Mean no-rate is our primary metric throughout this paper.

Inflection rate. Because no-rate is computed at every sentence boundary, it also produces a *curve* over the arc of the story. A complementary statistic is the **inflection rate**: the fraction of evaluated positions at which the smoothed no-rate curve sharply reverses direction. Because the vertex angle is a geometric measurement, it depends on the aspect ratio of the coordinate space. Before computing angles, we rescale each story’s curve onto a unit square: the x -axis (% story revealed) is mapped to $[0, 1]$ using the story’s position range, and the y -axis (smoothed no-rate) is independently mapped to $[0, 1]$ using the story’s own min and max. This fixes a 1:1 aspect ratio for angle computation, ensuring thresholds are comparable across stories of different lengths and no-rate ranges. We report results at three thresholds (30, 60, 120):

$$\text{inflection-rate}_\alpha(s) = \frac{|\{i : \theta_i \leq \alpha\}|}{|P_s|} \quad (3)$$

While mean no-rate summarizes overall unpredictability, inflection rate tracks narrative manipulation—the frequency with which a story reverses reader expectations. Figure 3 illustrates this with two *New Yorker* stories. Petrushevskaya’s “Poor Girl”, a realist story with a clinical and distanced exteriority, has high mean no-rate (0.925) but low inflection rate (0.048). “The Fellow” (Joy Williams), a surrealist story in which a guest at a writers’ residency vanishes in a flood and a dog begins to speak, has low mean no-rate (0.403) because the plot is minimal, but very high inflection rate (0.339): each small revelation briefly makes the ending guessable before the next twist displaces it. Whether these two statistics capture genuinely distinct dimensions of craft is an open question, but the contrast suggests that the no-rate curve contains information beyond its mean.

Post-spike retention. For each local peak in the no-rate curve (a position exceeding both neighbors), we measure how much of the peak value is retained at the minimum within the next 10 positions. The story-level **retention** is the mean across all peaks. Higher retention means tension peaks are sustained rather than immediately resolved.

Source	EQ-Bench	Group	Mean	Late	Retain	Infl ₆₀
GPT-5.2	83.80	<i>New Yorker</i>	0.765	0.607	52%	0.200
Claude Sonnet 4.6	82.04	Tell Me a Story	0.693	0.311	39%	0.128
Claude Opus 4.6	82.88	WritingPrompts	0.671	0.277	25%	0.084
<i>New Yorker</i> (professional)	78.71	StoryStar	0.656	0.375	35%	0.166
Tell Me a Story (amateur)	60.89	Top-10 LLMs	0.630	0.215	23%	0.106
StoryStar (amateur)	45.81					

Table 1: EQ-Bench rubric score (left) vs. our *100-Endings* metric (right). **Mean**: story-level no-rate; **Late**: mean no-rate at $\geq 80\%$ of story; **Retain**: % of peak no-rate retained after local peaks (higher = tension sustained longer); **Infl₆₀**: fraction of sharp ($\leq 60^\circ$) reversals. *New Yorker* leads on all four metrics while underperforming top LLMs on the EQ-Bench rubric.

3.2 Why a New Metric?

EQ-Bench Creative Writing v3 (Paech, 2023) is one of the most widely used benchmark for evaluating LLM-generated fiction. Its rubric scores stories on 22 criteria (e.g., adherence to instructions, prose quality, emotional complexity, etc.) using Claude Sonnet 4 as judge, producing a composite score on a 0–100 scale. Notably, none of the 22 directly evaluate for tension.²

We validate both EQ-Bench and our *100-Endings* metric on four groups of stories of varying quality (details see Appendix): *New Yorker* fiction (professional), Tell Me a Story (amateur, curated), StoryStar (amateur, uncurated), and LLM-generated stories by top performers from the EQ-Bench leaderboard (zero-shot outputs on 32 standard prompts).

Table 1 and Figure 8 report the results. On EQ-Bench, LLMs (81–84) comfortably outscore all human datasets (44–78), including the *New Yorker* stories. EQ-Bench’s LLM stories are produced in a single turn (with no system message, revision, or structural planning) from 32 genre-specific writing prompts, each 75–560 words specifying genre, character, setting, and stylistic constraints (e.g., a noir romance at Coney Island in 1937). *New Yorker* stories, by contrast, are written by professional authors, typically drafted and revised over months, and selected through a competitive editorial process at a magazine that has published literary fiction since 1925. It is implausible that a single-pass LLM could genuinely surpass this level of craft; if taken at face value, it would mean the craft of literary fiction has been solved—a claim directly contradicted by various human validation results from recent NLP research in this area (Marco et al., 2024; Chakrabarty et al., 2024; Tian et al., 2024; Marco et al., 2025; Ismayilzada et al., 2025). The more likely explanation is that the rubric misses something important, and EQ-Bench rewards dimensions where LLMs already excel (e.g., prose fluency, structural completeness, and instruction-following), while omitting narrative tension, the dimension where professional fiction most clearly surpasses LLM output.

Our *100-Endings* metric provides confirmation of our central hypothesis: LLMs cannot hold tension. *100-Endings* orders the five groups consistently with intuition: *New Yorker* (0.765) exceeds EQ-Bench’s top-10 LLMs (0.63) by +0.14, and all three amateur corpora fall in between (Table 1). The gap is especially stark in late-stage unpredictability: *New Yorker* stories maintain a late-stage no-rate of 0.607, nearly triple that of LLMs (0.215), which means that professional fiction preserves narrative tension even as it approaches its ending while LLM stories become predictable well before they conclude. Post-spike retention reinforces this: *New Yorker* stories retain 52% of local peaks on average, while LLMs retain only 23%. Professional authors raise narrative questions and leave them partially open; LLMs resolve them almost immediately. We argue that this is the missing dimension in EQ-Bench—the sustained management of what the reader does and does not yet know.

²While EQ-Bench also employs an Elo rating system based on pairwise matchups, we evaluate only the rubric scores. The leaderboard recently transitioned to Sonnet 4 as a judge (previously Sonnet-3.7). Examination of the repository indicates that while rubric scores have been updated, most Elo scores have not, making pairwise comparisons inconsistent. Furthermore, because the Elo ranking criteria closely mirror the initial rubric and entirely omit narrative tension, the same evaluative limitations apply to both.




Figure 4: Effect of our narrative pipeline on rubric scores and tension metrics (32 EQ-Bench prompts). Left: EQ-Bench rubric leaderboard with our pipeline at #1 (85.03). Right: 100-Ending on the same 32 prompts. The pipeline improves tension over vanilla models.

4 Pipeline

Story Idea Primer. The pipeline takes a story idea of any length as its sole input. These “seed” ideas can be user-supplied, drawn from existing benchmarks (e.g., the 32 genre-specific prompts from EQ-Bench that we use in Section 5), or produced by a custom LLM-based idea generation module that might combine genre profiles, tension-rich narrative tropes, or other mechanisms for combinatorial creativity to automatically produce diverse premises (see Appendix D.3).

Given a story idea, the pipeline (Figure 1) produces a full story in three sequential steps, threading a shared context through each stage so that every step can condition on all prior outputs.

Step 0: Warmup Analysis. We randomly select a reference story from one of two sources: full-text stories from our *New Yorker* dataset, or a curated pool of canonical short stories. For the latter, we provide only the title to leverage the model’s parametric knowledge of the story. An LLM (Claude Sonnet, $T=0.7$) reads (or recalls) the reference story and extracts 7 scene-level structural beats, each annotated with (1) the concrete action, (2) the specific narrative technique employed, and (3) what made it memorable. This step serves as a “warmup link,” priming the model with the structural logic of high-quality literary fiction before it encounters the target story idea.

Step 1: Beat Adaptation (The Thick Intermediate Layer). The extracted beat sheet template is adapted to the target story idea. The LLM produces a detailed intermediate plan, a structural “to-do list” that specifies the following for each beat: the narrative function, what information is revealed or withheld, the tension mechanism in play, and how stakes escalate. This intermediate representation is deliberately *thick*: it carries enough structural information to constrain the generated story away from default narrative templates while leaving stylistic-level decisions to the next step. When the warmup step is ablated, beats are generated from scratch, prompted with the story idea alone.

Step 2: Story Generation. The LLM generates the final story conditioned on the story idea and the thick intermediate plan from Step 1. The model is instructed to follow the structural to-do list closely, with a system-prompt emphasis on sustaining tension. When both planning steps are ablated, the pipeline reduces to single-shot (vanilla) generation, providing a controlled baseline.

5 Results

We investigate whether structured constraints can improve long-range narrative tension while maintaining or even enhancing other metrics of fiction quality. Despite the limitations of the EQ-Bench rubric demonstrated in Section 3, it remains one of the most widely used standards and captures dimensions of story quality orthogonal to tension. We apply our full three-step pipeline (Section 4)

Table 2: Pipeline effect on rubric quality and tension metrics (32 EQ-Bench prompts). The pipeline improves all four tension metrics across three models, with the largest gains on late-stage no-rate. For Sonnet 4.6, we run 4 independent repeats to validate the stability of *100-Endings* measurements (Appendix C.7).

Source	n	EQ-Bench	Mean NR	Late NR	Retain	IR ₆₀
Sonnet 4.6 (vanilla)	32	82.04	0.606	0.139	16.7%	0.077
Sonnet 4.6 + Pipeline	32	85.03	0.747	0.339	35.5%	0.117
Opus 4.6 (vanilla)	32	82.88	0.611	0.185	15.9%	0.090
Opus 4.6 + Pipeline	32	81.82	0.697	0.296	32.8%	0.142
GPT-5.2 (vanilla)	32	83.80	0.694	0.329	36.1%	0.131
GPT-5.2 + Pipeline	32	81.24	0.740	0.443	43.6%	0.141
Top-10 LLMs (0-shot)	319	81.88	0.630	0.215	23.3%	0.106
<i>New Yorker</i>	261	78.70	0.765	0.607	51.7%	0.200




Figure 5: No-rate curve for Updike’s “Snowing in Greenwich Village.” Tension rises once the husband walks the guest home alone ($\sim 60\%$ story revealed) and stays elevated through the ambiguous ending.

to three frontier models (Claude Sonnet 4.6, Claude Opus 4.6, and GPT-5.2) across the standard 32 EQ-Bench prompts, and evaluate the outputs on both the EQ-Bench rubric and *100-Endings*.

Table 2 reports the results. The pipeline improves all four tension metrics for all three models. The effect is strongest for Sonnet 4.6: mean no-rate rises from 0.606 to 0.747, late-stage no-rate more than doubles, and post-spike retention jumps from 16.7% to 35.5%. Opus 4.6 and GPT-5.2 show similar patterns, with late-stage no-rate gains of +0.111 and +0.114, respectively. For Sonnet 4.6, the pipeline also improves the EQ-Bench rubric score, placing it at #1 on the leaderboard (Figure 4). For Opus 4.6 and GPT-5.2, rubric scores decrease only slightly (-1.1 and -2.6 points), though we note the rubric’s insensitivity to tension (Section 3). The key finding is that the pipeline consistently closes a large fraction of the gap between zero-shot LLM writing and professional fiction on tension metrics—particularly late-stage no-rate, where the pipeline lifts all three models well above the Top-10 LLM baseline (0.215) and toward *New Yorker* levels (0.607).

6 Discussion

6.1 Case Study: What is the Shape of Stories With Tension?

The central finding of this work is that zero-shot LLMs tend to collapse narrative tension through premature resolution, a shortcoming that a thick, structured generation pipeline can mitigate. To illustrate the problem and the remedy, we examine the tension curves of three romance stories that share a single source of suspense: the unspoken chemistry between a man and a woman alone together. One is from our *New Yorker* dataset, while the other two are generated by CLAUDE SONNET 4.6 from identical prompts—one zero-shot, and one using structured pipeline. Because the genre, subject,