




Figure 6: No-rate curves for two stories generated by Sonnet 4.6 from the same prompt. **Left:** Zero-shot. **Right:** Narrative pipeline (warmup + planning). The zero-shot curve collapses early, while the pipeline version sustains high unpredictability. See Appendix F for full story texts.

and emotional register are held constant, the resulting differences in the no-rate curves isolate how effectively each story manages the gap between possibility and resolution.

The New Yorker. In Updike’s “Snowing in Greenwich Village” (1956), a husband walks an attractive female guest home through the snow while his wife stays behind with a cold. The moment introduces immediate stakes: will the underlying attraction culminate in an affair, or will nothing happen? As Figure 5 illustrates, the no-rate spikes once the pair is alone and stays elevated throughout their ambiguous parting at her door. Although the affair never materializes, the narrative sustains the possibility of transgression into the final sentence. Updike maintains the tension to the very end by never allowing the narrative to prematurely settle; the story remains suspended between seduction and restraint, making ambiguity the focal point.

The zero-shot failure. Figure 6 (left) illustrates the typical failure mode of an unguided LLM attempting a comparable romance, which exhibits a decisive turning point at 64% progression. Up to that mark, the banter sustains uncertainty as the owner initially rebuffs the actor. However, the moment she invites him to the back room for tea and acknowledges his honesty, their mutual attraction is explicitly confirmed. With the romantic trajectory resolved, the no-rate drops to 0.29 and completely flatlines by the final sentences. Unlike Updike, the vanilla model prematurely forecloses the narrative pathway and collapses tension the very moment the characters are alone together.

The pipeline recovery. Figure 6 (right) shows our pipeline’s version of the same prompt. The interaction is driven entirely by subtext, as the characters part without ever acknowledging their underlying chemistry after exchanging only oblique vulnerabilities. This deliberate ambiguity culminates in the story’s conclusion, as the distracted owner corrects her ledger: “She sat with the mistake for a moment before she fixed it.” Here, *mistake* and *fixed* operate on two levels, referring simultaneously to a mundane accounting error and the unacted-upon romantic spark. Consequently, the no-rate remains above 0.85 through 85% of the story. Like Updike, our pipeline sustains the tension between attraction without resolving into either acceptance or rejection. This highlights that the zero-shot failure is not an issue of prose fluency (both LLM versions are fluent), but a fundamental inability to *defer closure*—a deficit our pipeline avoids by design through thick planning.

6.2 What 100-Endings Measures

100-Endings is a measure of predictability, which we argue is a necessary (but not sufficient) component of tension. A story can be unpredictable and incoherent, since unpredictability without stakes, investment, or logic is simply chaos. The metric should therefore be used alongside mainstream benchmarks of generic writing quality. Yet, the metric’s ability to distinguish professional fiction from LLM outputs (unlike EQ-Bench) suggests it captures an essential element of narrative craft: the

deliberate management of reader expectations through information asymmetry, delayed resolution, and structural surprise.

While our four statistics (Section 3) primarily evaluate *plot-level* predictability, the generated endings offer a rich artifact for further analysis. Future work could develop metrics analyzing the *style* of these endings to capture sub-event forms of tension, such as thematic irony and psychological interiority.

The metric’s reliance on LLMs as generator raises a circularity concern: are we measuring the properties of the stories, or the idiosyncratic preferences of the judge? We mitigate this by strictly separating the generator and evaluator models (employing QWEN3 exclusively for evaluation) and by validating that the metric correctly orders human corpora of known quality (details see Appendix A).

The Gap to Professional Fiction Our best pipeline configuration achieves a mean no-rate of 0.702 on the EQ-Bench prompts, closing roughly half the gap between vanilla Sonnet 4.6 (0.577) and *New Yorker* fiction (0.765), but still below. The remaining gap likely reflects aspects of professional craft that our pipeline does not yet address: tacit knowledge based on cultural contexts, subversion of genre through deep familiarity with their conventions, and the kind of sentence-level surprise that emerges from a distinctive authorial voice rather than structural planning.

7 Conclusion and Future Work

Zero-shot LLMs struggle to hold narrative tension, resolving uncertainty prematurely rather than sustaining the gap between reader knowledge and narrative withholding. Because tension is downstream of hidden information—secrets, delayed revelations, chilling endings—this tendency to prematurely reveal produces stories that are polished but flat. The *100-Endings* metric makes this deficit measurable, operationalizing tension as unpredictability of endings at the sentence level. Our results reveal that stories generated by top LLMs exhibit the lowest tension of any evaluated corpus, even when rubric-based judges rate them above professional fiction. However, our structured generation pipeline demonstrates that this deficit is not immutable: literary theory-driven constraints can significantly improve long-range tension. While tension is a vital aspect of the craft of storytelling, it is only one element of the broader space of key narrative attributes required to build true narrative machines.

This paper provides a framework for quantifying and improving narrative tension, and represents an initial step toward a comprehensive computational narratology (Piper et al., 2021) that opens up several interesting directions: our immediate subsequent work will expand on various components of our pipeline, particularly the generation of synthetic story primers at scale; we also plan to conduct systematic ablations across every stage of the pipeline to isolate which structural constraints contribute most to tension—specifically, whether the key factor is the thickness of intermediate steps or the combinatorial novelty of the constraints themselves; aside from measuring tension, the *100-Endings* framework also produces a rich object (100 candidate endings at every sentence boundary) from which further metrics could be derived, particularly literary theory-driven measures of stylistic or thematic aspects of quality that operate below the level of overall tension.

Beyond the pipeline, tension is only one dimension of narrative craft. Good stories do not emerge from a vacuum, and our future work aims to explore the tacit knowledge and contextual frameworks behind narrative artifacts like tension—if multi-agent architectures can simulate the editorial processes, social dynamics, and cultural ecosystems from which high-quality literature originates. Modeling these environments may prove essential for closing the remaining gap to professional fiction.

Humanistic Mission Statement³

This work aims to broaden the intersection of AI research and the humanities, specifically literary theory. As LLMs become more deeply embedded in society, NLP research is growing increasingly interdisciplinary, encountering nuanced problem spaces that prove to be recalcitrant to the expertise and methods of computer science alone. We argue that it is vital to bridge NLP research with existing domain knowledge in relevant disciplines, such as applying narrative theory to computational

³Following Sui et al. (2025), we include this statement to ensure that our AI research faithfully adheres to the standards and values of humanities scholarship.

narrative understanding (Piper et al., 2021). Just as AI research has long embraced physics-informed machine learning to ground models in the natural world, we advocate for the development of literary-theory-informed models for computational creativity. We hope to encourage more collaborative work between NLP researchers and humanities scholars, ensuring that future models are grounded in the rigorous theoretical frameworks of the domains they attempt to emulate.

We strongly oppose the use of AI research to disrupt the craft of storytelling or even displace human writers—for instance, by building LLM systems that claim to write "better" stories in some shape or form. Our goal is to develop computational approaches that can study narrative and its affective powers at scale. Likewise, our *100-Endings* metric is not designed to reduce the quality of fiction to an algorithmic monolith, but to introduce a reliable and scalable yardstick to help us better understand the complex interplay between various mechanisms of storytelling.

Ethics Statement

Our research involves no human subjects and utilizes publicly available or archived literary datasets. To ensure ethical compliance, the *New Yorker* data is stored in a secure repository at our university; its use in this study constitutes fair use, as it is employed strictly for non-disseminating, non-commercial research. Other amateur corpora, such as StoryStar, were sourced from public-facing websites intended for open readership.

Impact on Human Creators. Utilizing computational tools to operationalize narrative tension allows us to map the formal contours of what is often felt as a nebulous spark of creativity. We hope this work has a positive impact on every person who takes satisfaction in a story finding its shape. Our objective is to demystify these structures to empower, rather than displace, the author. As stated in our humanistic mission statement, we do not intend for these systems to replace human storytellers; instead, we hope that by improving machine performance on these tasks, we create tools that augment the human imagination.

Acknowledgments

This work was supported, in part, by Schmidt Sciences’ Humanities and AI Virtual Institute (HAVI) and the “Humanistic AI” grant from the Neubauer Collegium for Culture and Society, University of Chicago. We thank the Textual Optics Lab at the University of Chicago for providing access to *The New Yorker* corpus.

References

- Raphaël Baroni. *La tension narrative: suspens, curiosité et surprise*. Seuil, 2007.
- William F Brewer and Edward H Lichtenstein. Stories are to entertain: A structural-affect theory of stories. *Journal of Pragmatics*, 6(5-6):473–486, 1982.
- Noël Carroll. *The Philosophy of Horror: Or, Paradoxes of the Heart*. Routledge, 1990.
- Tuhin Chakrabarty, Philippe Laban, Divyansh Agarwal, Smaranda Muresan, and Chien-Sheng Wu. Art or artifice? Large language models and the false promise of creativity. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*, 2024. URL <https://dl.acm.org/doi/10.1145/3613904.3642731>.
- Tuhin Chakrabarty, Philippe Laban, and Chien-Sheng Wu. Can AI writing be salvaged? Mitigating idiosyncrasies and improving human-AI alignment in the writing process through edits. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, 2025. URL <https://dl.acm.org/doi/full/10.1145/3706598.3713559>.
- Cyril Chhun, Pierre Colombo, Fabian M. Suchanek, and Chloé Clavel. Of human criteria and automatic metrics: A benchmark of the evaluation of story generation. In *Proceedings of the 29th International Conference on Computational Linguistics (COLING)*, pp. 5794–5836. International Committee on Computational Linguistics, 2022.

- Jeffrey C Ely, Alexander Frankel, and Emir Kamenica. Suspense and surprise. *Journal of Political Economy*, 123(1):215–260, 2015.
- Benjamin Feuer, Micah Goldblum, Teresa Datta, Sanjana Nambiar, Raz Besaleli, Samuel Dooley, Max Cembalest, and John P. Dickerson. Style outweighs substance: Failure modes of LLM judges in alignment benchmarking. In *The Thirteenth International Conference on Learning Representations*, 2025.
- G rard Genette. *Narrative Discourse: An Essay in Method*. Cornell University Press, 1980.
- Jian Guan and Minlie Huang. UNION: An unreferenced metric for evaluating open-ended story generation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 9157–9166. Association for Computational Linguistics, 2020.
- Alexander Gurung and Mirella Lapata. Learning to reason for long-form story generation. *arXiv preprint arXiv:2503.22828*, 2025. URL <https://arxiv.org/abs/2503.22828>.
- Fantine Huot, Reinald Kim Amplayo, Jennimaria Palomaki, Alice Shoshana Jakobovits, Elizabeth Clark, and Mirella Lapata. Agents’ room: Narrative generation through multi-step collaboration. In *Proceedings of the 2025 International Conference on Learning Representations*, 2025.
- Mete Ismayilzada, Claire Stevenson, and Lonneke van der Plas. Evaluating creative short story generation in humans and large language models. In *Proceedings of the 16th International Conference on Computational Creativity (ICCC’25)*. Association for Computational Creativity, 2025.
- Seungone Kim, Jamin Shin, Yejin Cho, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Seongjin Shin, Sungdong Kim, James Thorne, and Minjoon Seo. Prometheus: Inducing fine-grained evaluation capability in language models. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=8euJaTveKw>.
- Moritz Lehne and Stefan Koelsch. Toward a general psychological model of tension and suspense. *Frontiers in Psychology*, 6:79, 2015.
- Moritz Lehne, Pascal Engel, Martin Rohrmeier, Winfried Menninghaus, Arthur M Jacobs, and Stefan Koelsch. Reading a suspenseful literary text activates brain areas related to social cognition and predictive inference. *PLoS ONE*, 10(5):e0124550, 2015.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-eval: NLG evaluation using Gpt-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 2511–2522. Association for Computational Linguistics, 2023.
- Guillermo Marco, Julio Gonzalo, M. Teresa Mateo-Girona, and Ram n Del Castillo Santos. Pron vs prompt: Can large language models already challenge a world-class fiction author at creative text writing? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 19654–19670, Miami, Florida, USA, 2024. Association for Computational Linguistics.
- Guillermo Marco, Julio Gonzalo, and V ctor Fresno. The reader is the metric: How textual features and reader profiles explain conflicting evaluations of AI creative writing. In *Findings of the Association for Computational Linguistics: ACL 2025*, pp. 25432–25449, Vienna, Austria, 2025. Association for Computational Linguistics.
- Glenn Matlin, Devin Zhang, Rodrigo Barroso Loza, Diana M. Popescu, Joni Isbell, Chandreyi Chakraborty, and Mark Riedl. Do language models agree with human perceptions of suspense in stories?, 2025. URL <https://arxiv.org/abs/2508.15794>.
- Brian O’Neill and Mark O. Riedl. Dramatis: A computational model of suspense. In *Proceedings of the Twenty-Eighth AAAI Conference on Artificial Intelligence*, volume 28, pp. 944–950. AAAI Press, 2014. doi: 10.1609/aaai.v28i1.8836. URL <https://ojs.aaai.org/index.php/AAAI/article/view/8836>.