

- Vishakh Padmakumar and He He. Does writing with language models reduce content diversity? In *Proceedings of the 2024 International Conference on Learning Representations*, 2024.
- Samuel J. Paech. Eq-bench: An emotional intelligence benchmark for large language models. *arXiv*, 2023. doi: 10.48550/arxiv.2312.06281.
- Samuel J Paech. Eq-bench creative writing benchmark v3. <https://github.com/EQ-bench/creative-writing-bench>, 2025.
- Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. *arXiv preprint*, 2024.
- Chau Minh Pham, Jenna Russell, Dzung Pham, and Mohit Iyyer. Frankentext: Stitching random text fragments into long-form narratives. *arXiv preprint arXiv:2505.18128*, 2025. URL <https://arxiv.org/abs/2505.18128>.
- Andrew Piper, Richard Jean So, and David Bamman. Narrative theory for computational narrative understanding. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pp. 298–311, Online and Punta Cana, Dominican Republic, 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-main.26. URL <https://aclanthology.org/2021.emnlp-main.26/>.
- Vladimir Propp. *Morphology of the Folktale*. University of Texas Press, 1928.
- Meir Sternberg. Universals of narrative and their cognitivist fortunes (i). *Poetics Today*, 24(2): 297–395, 2003.
- Peiqi Sui. LLMs exhibit significantly lower uncertainty in creative writing than professional writers. *arXiv preprint arXiv:2602.16162*, 2026. URL <https://arxiv.org/abs/2602.16162>.
- Peiqi Sui, Juan Diego Rodriguez, Philippe Laban, J. Dean Murphy, Joseph P. Dexter, Richard Jean So, Samuel Baker, and Pramit Chaudhuri. KRISTEVA: Close reading as a novel task for benchmarking interpretive reasoning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 32829–32849, Vienna, Austria, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.acl-long.1577. URL <https://aclanthology.org/2025.acl-long.1577/>.
- Yufei Tian, Bodhisattwa Prasad Majumder, and Yulia Tsvetkov. Are large language models capable of generating human-level narratives? In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024. URL <https://aclanthology.org/2024.emnlp-main.978/>.
- Vera Tobin. *Elements of surprise: our mental limits and the satisfactions of plot*. Harvard University Press, Cambridge, Massachusetts, 2018. ISBN 978-0-674-98020-4. OCLC: 1002829918.
- Vivek Verma, Eve Fleisig, Nicholas Tomlin, and Dan Klein. Ghostbuster: Detecting text ghostwritten by large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pp. 1702–1717, Mexico City, Mexico, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.naacl-long.95. URL <https://aclanthology.org/2024.naacl-long.95/>.
- Lei Wang, Jianxun Lian, Yi Huang, Yanqi Dai, Haoxuan Li, Xu Chen, Xing Xie, and Ji-Rong Wen. CharacterBox: Evaluating the role-playing capabilities of LLMs in text-based virtual worlds. In *Proceedings of the 2025 Annual Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics (NAACL)*, pp. 6372–6391. Association for Computational Linguistics, 2025. URL <https://aclanthology.org/2025.naacl-long.323>.
- David Wilmot and Frank Keller. Modelling suspense in short stories as uncertainty reduction over neural representation. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pp. 1763–1788, Online, 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.acl-main.161. URL <https://aclanthology.org/2020.acl-main.161/>.

- Kaige Xie and Mark Riedl. Creating suspenseful stories: Iterative planning with large language models. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 2391–2407, St. Julian’s, Malta, 2024. Association for Computational Linguistics. URL <https://aclanthology.org/2024.eacl-long.147>.
- Weijia Xu, Nebojsa Jojic, Sudha Rao, Chris Brockett, and Bill Dolan. Echoes in ai: Quantifying lack of plot diversity in LLM outputs. *Proceedings of the National Academy of Sciences*, 122(35): e2504966122, 2025. doi: 10.1073/pnas.2504966122. URL <https://www.pnas.org/doi/10.1073/pnas.2504966122>.
- Qiang Yi, Yangfan He, Jianhui Wang, Xinyuan Song, ShiYao Qian, Xinhang Yuan, Yi Xin, Yijin Wang, Jingqun Tang, Yuchen Li, Junjiang Lin, Hongyang He, Zhen Tian, Tianxiang Xu, Keqin Li, Kuan Lu, Menghao Huo, Jiaqi Chen, Miao Zhang, Tianyu Shi, and Jianyuan Ni. SCORE: Story coherence and retrieval enhancement for AI narratives. *arXiv preprint arXiv:2503.23512*, 2025.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In *Advances in Neural Information Processing Systems*, volume 36, pp. 46595–46623. Curran Associates, Inc., 2023.

A Datasets

We evaluate on five story corpora spanning professional fiction, amateur writing, and LLM-generated text. Table 3 summarizes the key statistics.

Table 3: Summary of evaluation corpora. “100-Ending n ” is the number of stories used for the 100-ending metric (500–5,000 word range); “EQ-Bench n ” is the number used for EQ-Bench rubric evaluation (2,000–3,000 word range).

Corpus	Quality tier	Total stories	100-Ending n	EQ-Bench n
New Yorker	Professional	261	261	53
Tell Me a Story	Amateur (curated)	230	100	41
WP Human	Amateur (curated)	1,000	100	—
StoryStar	Amateur (scraped)	100	100	31
EQ-Bench LLMs	LLM-generated	$32 \times M$	319	319

New Yorker ($n = 261$). We utilize a corpus of short fiction published in *The New Yorker* between 1945 and 2019, widely regarded as a gold standard for contemporary literary fiction. To ensure ethical compliance, the data is stored in a secure repository at our university; its use in this study constitutes fair use, as it is employed strictly for non-disseminating, non-commercial research. From the total dataset, we select 261 stories under 5,000 words for the *100-Ending* analysis. For the EQ-Bench rubric evaluation, which requires a narrow word-count band to control for length bias (Section B), 53 stories fall within the 2,000–3,000 word range.

Tell Me a Story ($n_{\text{sample}} = 100$). Tell Me a Story (TMAS) is a publicly available dataset of short stories written by amateur authors in a guided creative-writing setting (Huot et al., 2025). The full corpus contains 230 stories. We randomly sample 100 stories (500–5,000 word range) for the 100-ending evaluation. For the EQ-Bench rubric evaluation, 41 stories fall within the 2,000–3,000 word band.

WP Human ($n_{\text{sample}} = 100$). We draw human-written stories from the *WritingPrompts* subset of the Ghostbuster dataset (Verma et al., 2024), which collects Reddit `r/WritingPrompts` submissions paired with their prompt. The full subset contains 1,000 stories. We randomly sample 100 for the 100-ending evaluation. Because all stories in this corpus fall below 1,800 words, none meet the 2,000-word minimum for EQ-Bench rubric evaluation; WP Human is therefore excluded from the rubric analysis.

StoryStar ($n = 100$). StoryStar is a corpus of 100 amateur short stories scraped from `storystar.com`, a public creative-writing platform. Stories are sampled to fall within the 500–5,000 word range. Unlike TMAS, StoryStar stories receive no editorial curation, making this corpus a lower-quality amateur baseline.

EQ-Bench LLMs ($n = 320$). We include zero-shot LLM-generated stories from the top-performing models on the EQ-Bench Creative Writing leaderboard (Paech, 2025). Each model generates one story per each of the 32 standard EQ-Bench prompts, sourced directly from the benchmark’s public repository.⁴ We use outputs from the top-10 ranked models at the time of evaluation. After filtering to the 500–5,000 word range, this yields 319 stories. These serve as the baseline: stories that score highly on existing rubric-based metrics but, as we show, lack narrative tension.

B The EQ Bench Creative Writing Baseline

EQ-Bench has tension-adjacent criteria, but they are insufficient. The rubric is not entirely blind to tension. Among its 22 criteria are several that relate to engagement and surprise: *Emotionally Engaging*, *Emotionally Complex*, *Overall Reader Engagement*, *Well-earned Lightness/Darkness*, and

⁴<https://github.com/EQ-bench/creative-writing-bench>




Figure 7: Tension-related EQ-Bench sub-criteria across three human-written corpora (higher = better; negative criteria are inverted). The rubric captures some tension-related signal but evaluates it as a holistic *judgment*—not through the operational lens of how a story manages information and prediction over time.

(inverted) *Unsurprising / Uncreative*, *Meandering*, *Incongruent Ending Positivity*, and *Unearned Transformations*. Figure 7 shows that these criteria do differentiate the three human corpora in the expected direction: *New Yorker* fiction substantially outperforms amateur writing on every tension-related sub-score. Obviously, these criteria can indeed distinguish good writing from bad writing.

Yet these criteria evaluate tension as a *holistic rubric judgment*: a judge reads the complete story and rates how “surprising” or “engaging” it felt. This approach has two structural limitations. First, rubric criteria like “Unsurprising / Uncreative” could conflate multiple constructs—an inventive premise can score well even if the plot unfolds predictably. Second, a single post-hoc score cannot capture *where* and *how* tension operates across the story arc: whether information is withheld, whether resolution is delayed, or whether the narrative resists prediction at each point in its progression. As a result, EQ Bench cannot distinguish great stories from ones that just have good prose. As our formal definitions in Section 2 make clear, measuring tension requires a **position-level** evaluation that tracks predictability across the entire arc.

Length-controlled evaluation. LLM judges exhibit a strong length bias in pairwise and rubric-based evaluation: longer outputs systematically receive higher scores (Paech, 2025). EQ-Bench mitigates this by truncating outputs at 4,000 characters, placing all models on equal footing. We adopt an analogous strategy for corpus evaluation: we restrict the rubric assessment to stories in a narrow word-count band (2,000–3,000 words), ensuring that differences in rubric scores reflect writing quality rather than length. This yields 31 StoryStar stories, 41 TMA stories, and 53 New Yorker stories for the human corpora.

We exclude *WritingPrompts* (WP Human) stories from the rubric evaluation because all stories in the corpus fall below 1,800 words, with zero stories reaching the 2,000-word minimum. The WP Human corpus is still included in the 100-ending analysis (Section C), where the metric operates at the sentence level and is robust to length variation (Appendix E).

For LLM-generated stories, most zero-shot outputs from the top-10 EQ-Bench models are substantially shorter than 2,000 words: across 319 stories (10 models $\times$ 32 prompts), the median length is 1,102 words, with only 14% falling in the 2,000–3,000 word range. GPT-5.2 is the exception, with 23 of its 32 stories in the target range. We therefore report two LLM entries: GPT-5.2 restricted to 2,000–3,000 words (the length-controlled comparison), and the pooled top-10 at all lengths. Note that the length restriction *disadvantages* LLMs: EQ-Bench’s own documentation notes that shorter outputs receive lower rubric scores. The fact that even at all lengths—where LLMs are penalized by their own brevity—they still outscore all human corpora makes the dissociation with our 100-ending metric all the more striking.