

Table 4: EQ-Bench Creative Writing rubric scores (0–100 scale). **Overall**: mean across all 22 criteria. **Tension₈**: mean across the 8 tension-related criteria only (*Emotionally Engaging*, *Emotionally Complex*, *Overall Reader Engagement*, *Well-earned Lightness/Darkness*, and inverted *Unsurprising/Uncreative*, *Meandering*, *Incongruent Ending Positivity*, *Uneared Transformations*). Human corpora restricted to 2,000–3,000 words; GPT-5.2 similarly restricted; pooled top-10 includes all lengths. Judge: Claude Sonnet 4.

Source	Overall	Tension ₈	±95% CI	n
GPT-5.2 (2–3k words)	84.28	83.86	1.81	23
Top-10 LLMs (all lengths)	81.88	80.80	0.55	319
<i>New Yorker</i> (professional)	78.61	78.40	4.33	53
Tell Me a Story (amateur)	60.89	55.85	3.05	41
StoryStar (amateur)	45.81	44.17	4.23	31

Full rubric results. Table 4 reports the EQ-Bench rubric scores across all evaluated groups. All evaluations use Claude Sonnet 4 as judge, following the standard EQ-Bench protocol (Paech, 2025).

GPT-5.2, even when restricted to the same 2,000–3,000 word range as the human corpora, scores 84.28 overall—comfortably above *New Yorker* fiction (78.70). The pooled top-10 at all lengths scores 81.88, still above professional fiction despite the length disadvantage. Notably, even on the 8 tension-related sub-criteria alone, LLMs outscore *New Yorker*: GPT-5.2 scores 83.86 vs. 78.40 for professional fiction. This means the rubric’s tension-adjacent criteria—which evaluate tension as a holistic post-hoc judgment—do not capture the same signal as our position-level 100-ending metric, which correctly reverses the ordering. Among human corpora, *New Yorker* leads by a wide margin over Tell Me a Story (60.89 overall, 55.85 on tension) and StoryStar (45.81 overall, 44.17 on tension).

C The 100-Ending Metric

This section formalizes the 100-ending metric and provides complete implementation details, including exact prompt templates and all hyperparameters.

C.1 Overview

The 100-ending metric measures **narrative predictability** at the sentence level. The intuition is simple: at each point in a story, we ask a language model to predict how the story ends, repeating this 100 times, then check how many predictions match what actually happens. If the true ending is easy to predict, most predictions will match; if it is surprising, few will.

Crucially, the model is asked to *describe* how the story ends—a plot-level prediction—rather than to write creative prose. This frames the task as forecasting rather than generation, measuring whether the model can anticipate the narrative trajectory.

Formally, let $S = (s_1, s_2, \dots, s_N)$ be a story decomposed into N sentences. At each sentence boundary $i \in \{1, \dots, N-1\}$, we:

1. **Predict** 100 independent ending predictions $\hat{e}_1^{(i)}, \dots, \hat{e}_{100}^{(i)}$ conditioned on the prefix $(s_1, \dots, s_i)$, using a generation model $\mathcal{G}$.
2. **Judge** each prediction against the true remainder $(s_{i+1}, \dots, s_N)$ using a judge model $\mathcal{J}$, which returns a binary verdict $j_k^{(i)} \in \{0, 1\}$ indicating whether $\hat{e}_k^{(i)}$ captures the “general direction” of the true continuation.
3. **Aggregate** into a position-level **no-rate**:

$$\text{no-rate}(i) = 1 - \frac{1}{100} \sum_{k=1}^{100} j_k^{(i)} \quad (4)$$

A no-rate of 0.85 means 85 of 100 predicted endings fail to match what actually happens—the story is highly unpredictable at that point. The **story-level no-rate** is the mean across all evaluated positions:

$$\overline{\text{no-rate}} = \frac{1}{|\mathcal{P}|} \sum_{i \in \mathcal{P}} \text{no-rate}(i) \quad (5)$$

where $\mathcal{P}$ is the set of retained positions (see Section C.3).

C.2 Sentence Splitting

Stories are split into sentences using a rule-based boundary detector. The algorithm identifies boundaries where terminal punctuation (., !, ?), optionally followed by closing quotation marks or parentheses, is followed by whitespace and then an uppercase letter or opening punctuation. The regex pattern is:

```
[.!?] [ "\"\u2019) ] *\s+ (?= [A-Z\u201c" ( \[ ] )
```

This heuristic correctly handles most English prose, including quoted dialogue and parenthetical asides. Both the generation and judge pipelines use an identical implementation of this splitter.

C.3 Position Selection

Not all sentence boundaries are evaluated. We apply two filters:

1. **Minimum context:** Position i is skipped if the preceding sentence s_{i-1} contains fewer than 10 words (`MIN_SENT_WORDS = 10`). This excludes positions following very short sentences (e.g., single-line dialogue tags) that provide minimal narrative context.
2. **Token percentage window:** We compute the fraction of story tokens revealed at each position using the tokenizer of the generation model. Only positions where 10%–99% of tokens have been revealed are retained. The lower bound excludes positions with insufficient context; the upper bound excludes trivial positions near the final sentence.

Stories exceeding 5,000 words are excluded entirely to ensure the full context fits within the generation model’s context window at every position.

C.4 Stage 1: Ending Generation

Model. We use **Qwen3-32B** as the generation model $\mathcal{G}$, served via vLLM with offline batch inference on a single NVIDIA H200 GPU. The vLLM engine is configured with `dtype=bfloat16, max_model_len=8192, enable_prefix_caching=True`, and `gpu_memory_utilization=0.92`. The model’s chain-of-thought mode is disabled (`enable_thinking=False`).

Prompt template. At each retained position i , we construct a chat-formatted prompt with a system message and user message:

```
System:
You are a fiction writer. Write only story text.

User:
In 2–3 sentences, describe how you think this story ends.
No commentary, no explanation, no preamble.
STORY SO FAR:
{s1 || s2 || ... || si}
```

where `||` denotes string concatenation with a space separator. The prompt asks the model to *describe* the ending rather than write it as prose. This frames the task as plot-level forecasting: the model predicts what events occur, not how they would be narrated.




Figure 8: Overall (left) and late-stage ($\geq 80\%$, right) mean no-rate by corpus. *New Yorker* fiction is the least predictable on both measures. Top-10 LLMs score lowest—below all human corpora, including amateurs—despite outscoring *New Yorker* on the EQ-Bench rubric (Table 4).




Figure 9: Inflection rate by corpus at three vertex-angle thresholds (30, 60, 120). *New Yorker* fiction has the highest inflection rate at all thresholds, indicating more frequent sharp reversals in the no-rate curve. The gap narrows at looser thresholds as all corpora accumulate more qualifying reversals.




Figure 10: Post-spike tension retention (%) by corpus. After a local peak in the no-rate curve, what fraction of the peak value is retained at the subsequent minimum? *New Yorker* stories retain the most tension after peaks (median 54%), while LLM stories resolve peaks almost completely (median 22%).

Table 5: Ending generation sampling parameters.

Parameter	Value	Rationale
Temperature	1.2	Elevated to produce diverse predictions that probe the full space of plausible endings. Lower temperatures collapse the prediction distribution, making it impossible to distinguish “predictable because all readers would agree” from “predictable because the model happens to be right.”
Top- p (nucleus)	0.95	Standard nucleus sampling threshold.
Max tokens	150	Sufficient for 2–3 sentences of plot-level description.
n (samples per prompt)	100	Each position yields exactly 100 independent predictions.
Batch size	64	Positions processed in batches of 64 for GPU utilization.

Sampling parameters. We generate $K = 100$ independent ending predictions per position using vLLM’s `n=100` parameter (all 100 samples drawn in a single batched call per position). The sampling configuration is:

Design rationale: temperature. The choice of $T = 1.2$ is deliberate. At $T = 0$ (greedy), the model produces a single deterministic prediction—this collapses the metric to a binary pass/fail at each position. At $T = 1.0$, predictions are diverse but still cluster around high-probability narratives. At $T = 1.2$, the distribution is broad enough to probe the full space of plausible endings while remaining coherent. We verified empirically that increasing T beyond 1.2 produces degraded output quality without improving metric sensitivity.

C.5 Stage 2: LLM Judge

Model. We use **Qwen3-8B** as the judge model $\mathcal{J}$, served via vLLM with `dtype=bfloat16`, `max_model_len=16384`, and **temperature 0.0** (greedy decoding) to ensure deterministic, reproducible judgments. Maximum generation length is 10 tokens. Batch size for judge inference is 512 prompts per vLLM call. Chain-of-thought is disabled.

System prompt.

```
You are a literary judge. You determine whether a predicted
story ending captures the general direction of the true
story continuation. Answer Y if the prediction gets the
broad plot trajectory right -- the same characters are
involved, similar types of events unfold, or the emotional
arc matches. Minor differences in specific details,
wording, or sequence are acceptable. Answer N only if the
prediction describes a fundamentally different outcome from
what actually happens.
```

User prompt. Each of the 100 predicted endings at a given position is judged **independently** in its own prompt (i.e., the judge sees only one prediction at a time, not all 100). For position i and predicted ending $\hat{e}_k^{(i)}$:

```
Here is the actual continuation of a story:
--BEGIN CONTINUATION--
{si-1} {si || si+1 || ... || sN}
--END CONTINUATION--
Here is a predicted ending:
--BEGIN PREDICTED--
{si-1} { $\hat{e}_k^{(i)}$ }
--END PREDICTED--
Does the predicted ending capture the general direction
of the true continuation -- similar characters involved,
```