

similar types of events, or a matching emotional arc?
 Answer Y if the broad trajectory is right, even if specific
 details differ. Answer N only if the prediction is
 fundamentally wrong about what happens. Answer ONLY Y or
 N.
 Judgment:

Context anchoring. Both the actual continuation and the predicted ending are prefixed with the same **preceding sentence** s_{i-1} . This shared anchor provides context when either the actual continuation or the prediction is too brief or vague on its own. The actual continuation $(s_i, \dots, s_N)$ is truncated to at most 2,000 words to stay within context limits.

Verdict extraction. The judge’s output is parsed by checking whether the first character (after stripping whitespace and uppercasing) is Y. Any other response (including N, empty, or malformed) is counted as a rejection ($j_k = 0$).

Design rationale: independent judging. We judge each prediction independently rather than presenting all 100 in a single prompt for two reasons: (1) context length—100 predictions plus the actual continuation would exceed most context windows; (2) independence—presenting multiple predictions in a single prompt introduces ordering effects and cross-contamination between judgments.

C.6 Aggregation and Reporting

The 100-ending framework produces a position-level no-rate curve for each story. We derive four story-level statistics from this curve.

1. Mean no-rate. The primary metric. At each position i , the no-rate is the fraction of 100 predictions the judge rejects:

$$\text{no-rate}(i) = 1 - \frac{1}{100} \sum_{k=1}^{100} j_k^{(i)} \quad (6)$$

The story-level mean no-rate averages across all retained positions:

$$\overline{\text{no-rate}} = \frac{1}{|\mathcal{P}|} \sum_{i \in \mathcal{P}} \text{no-rate}(i), \quad \mathcal{P} = \{i : 10\% \leq \text{pct}(i) \leq 99\%\} \quad (7)$$

2. Late-stage no-rate. The mean no-rate restricted to positions where $\geq 80\%$ of the story has been revealed. This measures whether the story remains unpredictable as it approaches its ending, or whether the conclusion becomes obvious once most of the narrative has been read.

3. Post-spike convergence. For each local peak in the no-rate curve (a position where no-rate exceeds both neighbors), we compute the percentage drop from the peak to the minimum within the next 10 positions. The story-level metric is the mean of these drops across all peaks. More negative values indicate that tension peaks resolve more sharply—the story answers narrative questions decisively after raising them.

4. Inflection rate. The fraction of positions at which the smoothed no-rate curve sharply reverses direction, as described in Section 3. Before computing vertex angles, both axes are rescaled to $[0, 1]$ (using each story’s own range) to fix a 1:1 aspect ratio. We report at three thresholds: $\alpha \in \{30, 60, 120\}$.

Group-level reporting. For a group of n stories, we report the mean of story-level values. Confidence intervals are $\bar{x} \pm 1.96 \cdot \sigma / \sqrt{n}$.

Full results. Table 6 reports all metrics across the evaluated corpora.

New Yorker fiction leads on mean no-rate, late-stage no-rate, and all three inflection-rate thresholds. The gap is especially pronounced in late-stage no-rate (0.607 vs. 0.28–0.38 for amateur corpora),

Table 6: Full 100-ending metric results across corpora (stories filtered to 500–5,000 words). Mean NR = mean no-rate; Late NR = mean no-rate at $\geq 80\%$ of story; Retain = % of peak no-rate retained after local peaks (higher = tension sustained); IR_α = inflection rate at vertex angle $\leq \alpha$.

Corpus	n	Mean NR	Late NR	Retain	IR_{30}	IR_{60}	IR_{120}
<i>New Yorker</i> (professional)	261	0.765	0.607	51.7%	0.128	0.200	0.271
Tell Me a Story (amateur)	100	0.693	0.311	39.3%	0.058	0.128	0.229
WP Human (amateur)	100	0.671	0.277	25.2%	0.028	0.084	0.211
StoryStar (amateur)	100	0.656	0.375	34.7%	0.105	0.166	0.232
Top-10 LLMs (0-shot)	319	0.630	0.215	23.3%	0.049	0.106	0.201

indicating that professional fiction sustains unpredictability through the final act. Post-spike convergence shows the opposite pattern: *New Yorker* stories have the *shallowest* post-spike drops (-48.3%), meaning tension peaks are not fully resolved—consistent with the literary technique of leaving narrative questions partially open.

C.7 Measurement Stability

A natural concern is whether the stochastic generation step ($T=1.2$) introduces unacceptable run-to-run variance in the no-rate metric. We evaluate stability by running the full pipeline 4 independent times on the same 32 EQ-Bench stories for both vanilla and pipeline-generated Sonnet 4.6 stories. Each repeat generates 100 fresh endings at every sentence position and re-judges them.

Table 7: Measurement stability across 4 independent evaluation repeats ($T_{\text{gen}}=1.2$, $n=32$ iter-1 stories). Standard deviations across repeats are ≤ 0.001 on mean no-rate and ≤ 0.001 on late-stage no-rate, indicating negligible measurement noise.

Config	Rep	Mean NR	Late NR	IR_{60}	Retain
Sonnet 4.6 (vanilla)	Rep 1	0.606	0.137	0.073	14.6%
	Rep 2	0.606	0.139	0.077	16.7%
	Rep 3	0.607	0.137	0.066	13.5%
	Rep 4	0.605	0.137	0.067	12.9%
	Mean $\pm$ Std	0.606 ± 0.001	0.138 ± 0.001	0.071 ± 0.005	$14.4 \pm 1.4\%$
Sonnet 4.6 + Pipeline	Rep 1	0.749	0.340	0.135	34.1%
	Rep 2	0.747	0.340	0.124	35.5%
	Rep 3	0.747	0.339	0.117	35.5%
	Rep 4	0.747	0.337	0.123	36.7%
	Mean $\pm$ Std	0.747 ± 0.001	0.339 ± 0.001	0.125 ± 0.007	$35.4 \pm 0.9\%$

Result. Mean no-rate and late-stage no-rate are stable to ± 0.001 across repeats—well within the 95% CI of the cross-story mean (± 0.05 – 0.06). Inflection rate and retention show slightly more variance (± 0.005 – 0.007 and ± 0.9 – 1.4% , respectively), consistent with these being second-order statistics that amplify small position-level fluctuations. The pipeline effect (vanilla $\rightarrow$ pipeline) is consistent across all 4 repeats: mean NR improves by $+0.141 \pm 0.001$, late NR by $+0.201 \pm 0.002$. These results confirm that even at generation temperature $T=1.2$, the 100-ending metric produces highly reproducible measurements.

D Pipeline Details

D.1 Pipeline Architecture

The narrative generation pipeline consists of up to three stages, each independently ablatable:

1. **Step 0 — Warmup:** Extract a structural beat sheet from a reference story (either a famous title or a full New Yorker story text).

2. **Step 1 — Beat Adaptation:** Adapt the warmup beat sheet to the target story idea. When warmup is absent, generates beats from scratch.
3. **Step 2 — Story Writing:** Write the final story from the adapted beats, or directly from the idea in vanilla mode.

All API calls use the OpenRouter endpoint. Reasoning/thinking is disabled for all calls. Stories shorter than 200 characters are automatically discarded and regenerated.

D.2 Shared Hyperparameters

Table 8 lists the hyperparameters shared across all pipeline steps (except where noted).

Table 8: Shared generation hyperparameters for story pipeline steps.

Parameter	Value	Notes
Temperature	0.7	All pipeline steps
Max tokens	4,000	Per API call
min_p	0.1	Minimum probability sampling
Reasoning	Disabled	max_tokens:0 for reasoning
Parallel threads	4–6	Per configuration run

D.3 Pipeline Input

Our structured generation pipeline is designed to take an initial “seed” idea and autonomously handle all subsequent structural planning and narrative execution. This seed input can be highly flexible in length. The pipeline accommodates anything from a single-sentence premise to a detailed outline. For readers and researchers interested in utilizing or adapting this pipeline, seed ideas typically originate from one of three sources:

- **Existing Benchmark Prompts:** Prompts from established writing benchmarks, such as the EQ-Bench prompts utilized in our evaluation. They provide a baseline for comparative analysis.
- **User-Generated Ideas:** Human authors can supply their own story premises. While this allows for maximum creative control, the qualitative depth and structural viability of the user’s initial input will naturally impact the quality of the pipeline’s final output—a variable that remains difficult to control or standardize in evaluations.
- **Synthetically Generated Ideas:** Seed premises can also be generated by LLMs using structured, theory-driven constraints. For instance, to automate genre blending, an LLM can be prompted with a combinatorial profile sampled from the *Screenwriters Taxonomy* (e.g., macro-genre, voice, narrative pathway, and POV) alongside narrative building blocks curated from an extensive list of tropes (such as *TV Tropes*). This approach reliably produce structurally sound, 3–4 sentence “pitch”-like ideas that lean into unusual combinatorial spaces. While this approach can yield diverse premises, finding the optimal configuration for AI-driven seed generation remains an open question for future work.

D.4 Pipeline Prompts

D.4.1 Step 0: Warmup Prompts

Step 0 generates a structural beat sheet from a reference story. One reference is randomly selected per story generation. Two warmup modes are tested:

Title mode. The model receives only a famous story title and generates beats from its training knowledge. The reference pool consists of 3 canonical short stories:

1. *The Ones Who Walk Away from Omelas* — Ursula K. Le Guin
2. *The Lottery* — Shirley Jackson
3. *The Yellow Wallpaper* — Charlotte Perkins Gilman

For future work, we recommend using a longer list of 20 canonical short stories spanning diverse narrative techniques, such as:

1. *The Ones Who Walk Away from Omelas* — Ursula K. Le Guin
2. *The Lottery* — Shirley Jackson
3. *The Yellow Wallpaper* — Charlotte Perkins Gilman
4. *Hills Like White Elephants* — Ernest Hemingway
5. *A Good Man Is Hard to Find* — Flannery O'Connor
6. *The Garden of Forking Paths* — Jorge Luis Borges
7. *The Metamorphosis* — Franz Kafka
8. *Rashōmon* — Ryūnosuke Akutagawa
9. *The Necklace* — Guy de Maupassant
10. *The Overcoat* — Nikolai Gogol
11. *Girl* — Jamaica Kincaid
12. *Axolotl* — Julio Cortázar
13. *Cathedral* — Raymond Carver
14. *The Lady with the Dog* — Anton Chekhov
15. *Araby* — James Joyce
16. *A Very Old Man with Enormous Wings* — Gabriel García Márquez
17. *Sonny's Blues* — James Baldwin
18. *The Fifth Story* — Clarice Lispector
19. *A Madman's Diary* — Lu Xun
20. *The Bear Came Over the Mountain* — Alice Munro

The title-mode prompt (style `cl_specific`) is:

```
Using "{ref_name}" as your reference, extract 7 scene-level
beats. For each beat describe: (1) the concrete action,
(2) the specific narrative technique used, and (3) identify
what made it memorable. What narrative trick, thematic
resonance, or structural choice made each story work?
```

Full-text mode. The model reads the complete text of a New Yorker short story and extracts structural patterns. The reference pool consists of:

1. *The Fellow* — Joy Williams
2. *Poor Girl* — Ludmilla Petrushevskaya
3. *The Dog* — J. M. Coetzee

For future work, we recommend using a longer list of 20 New Yorker stories, such as:

1. *The Fellow* — Joy Williams
2. *Poor Girl* — Ludmilla Petrushevskaya
3. *The Dog* — J. M. Coetzee
4. *I Can Speak!*TM — George Saunders
5. *She Said He Said* — Hanif Kureishi
6. *Dandelion* — Lore Segal
7. *Flaubert Again* — Anne Carson
8. *Riddle* — Thomas McGuane
9. *The Hotel* — Anne Enright
10. *Everything Is Far from Here* — Cristina Henríquez
11. *The Adventure of a Skier* — Italo Calvino
12. *The Piano Teacher's Pupil* — William Trevor
13. *Fly Already* — Etgar Keret
14. *Solstice* — Anne Enright
15. *Crazy They Call Me* — Zadie Smith
16. *Ladies' Lunch* — Lore Segal
17. *The Hanging of the Schoolmarm* — Robert Coover
18. *Back the Way You Went* — Anne Carson
19. *Two Men Arrive in a Village* — Zadie Smith