

Section	Metric	Prompting	Constrained		Baseline
			a priori	online	no forb.
Eval	True accuracy	97.40	92.30	87.10	34.00
	% exact viol.	1.00	0.00	1.00	45.90
	% only morphol. viol.	0.50	0.00	0.50	17.50
	% no <eot>	1.00	2.10	2.10	0.00
	% target leaked	2.10	7.70	12.40	4.10
gpt-oss judge	pass@1	0.00	0.50	0.00	0.00
	pass@2	0.00	1.00	0.50	0.50
	pass@3	0.00	1.50	0.50	0.50
	pass@5	0.50	2.10	1.00	4.60
	pass@10	1.50	5.70	4.60	6.70
	min	1e-6	2e-6	2e-6	1e-6
	p25	1.0e-5	1.1e-5	1.0e-5	1.4e-5
	median	1.8e-5	2.7e-5	2.2e-5	3.8e-5
	avg	3.4e-4	3.5e-3	1.9e-3	2.3e-3
	p75	6.1e-5	1.6e-4	1.9e-4	2.5e-4
	max	0.0269	0.2779	0.1916	0.1281
gemma judge	pass@1	0.50	0.00	0.00	0.00
	pass@2	4.60	5.20	4.10	4.10
	pass@3	6.70	7.70	7.70	9.80
	pass@5	11.30	13.90	10.30	17.50
	pass@10	16.50	19.60	20.60	21.60
	min	0	0	0	0
	p25	0	0	0	0
	median	0	0	0	0
	avg	2.2e-3	7.0e-4	3.6e-4	4.7e-4
	p75	5e-6	8e-6	8e-6	9e-6
	max	0.2870	0.0746	0.0228	0.0309

Table 6
Evaluation results – **gpt-oss-20b**